

954:534 Wish Project

Kanya Kreprasertkul, Yelin Shin and Zhe Ren

```
options(warn = -1)
library(dplyr)
library(tidyr)
#library(tidyverse)
library(GGally)
library(plotly)
library(cowplot)
library(ggcorrplot)
library(stringr)
```

Data pre-processing

```
wish <- read.csv('summer-products-with-rating-and-performance_2020-08.csv')  
  
#dropping unnecessary columns  
drops <- c('title', 'tags', 'crawl_month', 'theme', 'product_id', 'product_picture', 'product_url', 'merchant_id', 'merchant_profile_picture', 'merchant_info_subtitle', 'merchant_name', 'merchant_title', 'urgency_text', 'title_orig', 'shipping_option_name', 'currency_buyer')  
wish <- wish[, !(names(wish) %in% drops)]  
  
#convert NA to 0  
wish$has_urgency_banner <- as.integer(wish$has_urgency_banner)  
wish$has_urgency_banner[which(is.na(wish$has_urgency_banner))] <- 0  
wish$rating_five_count[which(is.na(wish$rating_five_count))] <- 0  
wish$rating_four_count[which(is.na(wish$rating_four_count))] <- 0  
wish$rating_three_count[which(is.na(wish$rating_three_count))] <- 0  
wish$rating_two_count[which(is.na(wish$rating_two_count))] <- 0  
wish$rating_one_count[which(is.na(wish$rating_one_count))] <- 0  
wish$rating[which(wish$rating_count == 0)] <- 0  
  
# cleaning size and color option  
wish <- wish %>%  
  mutate(product_variation_size_id = tolower(product_variation_size_id)) %>%  
  mutate(product_variation_size_id = gsub(pattern = '.', replacement = '',  
                                         x = product_variation_size_id, fixed = TRUE)) %>%  
  mutate(product_variation_size_id = gsub(pattern = '(size-*|)(size)', replacement = '',  
                                         x = product_variation_size_id)) %>%  
  mutate(product_variation_size_id = gsub(pattern = '.+[-]', replacement = ''
```

```

mutate(product_variation_size_id = ifelse(grepl(pattern = 'xl', product_variation_size_id),
                                           'xl', product_variation_size_id))
%>%
mutate(product_variation_size_id = ifelse(grepl(pattern = 'xs', product_variation_size_id),
                                           'xs', product_variation_size_id))
%>%
mutate(product_variation_size_id = str_replace(product_variation_size_id, '
', '')) %>%
mutate(product_variation_size_id = ifelse(product_variation_size_id %in% c(
's', 'xs', 'm', 'l', 'xl'), product_variation_size_id, 'One-sized'))
wish <- wish %>%
  mutate(product_color = tolower(product_color)) %>%
  mutate(product_color = ifelse(grepl(pattern = 'red|burgundy|claret|wine|jasper', product_color),
                                'red', product_color)) %>%
  mutate(product_color = ifelse(grepl(pattern = 'blue|navy', product_color),
                                'blue', product_color)) %>%
  mutate(product_color = ifelse(grepl(pattern = 'white', product_color),
                                'white', product_color)) %>%
  mutate(product_color = ifelse(grepl(pattern = 'green|army', product_color),
                                'green', product_color)) %>%
  mutate(product_color = ifelse(grepl(pattern = 'black', product_color),
                                'black', product_color)) %>%
  mutate(product_color = ifelse(grepl(pattern = 'yellow|leopard|gold', product_color),
                                'yellow', product_color)) %>%
  mutate(product_color = ifelse(grepl(pattern = 'pink|rose', product_color),
                                'pink', product_color)) %>%
  mutate(product_color = ifelse(grepl(pattern = 'grey|gray|silver', product_color),
                                'gray', product_color)) %>%
  mutate(product_color = ifelse(grepl(pattern = 'purple|violet', product_color),
                                'purple', product_color)) %>%
  mutate(product_color = ifelse(grepl(pattern = 'orange|apricot', product_color),
                                'orange', product_color)) %>%
  mutate(product_color = ifelse(grepl(pattern = 'beige|nude|ivory|coffee|brown|khaki|camel',
                                product_color), 'khaki', product_color)) %>%
  mutate(product_color = ifelse(grepl(pattern = 'floral|multicolor|camouflage|rainbow|star',
                                product_color), 'multicolor', product_color))

```

```

#name blank category
wish['product_color'][wish['product_color'] == ''] <- 'Not defined'
wish['origin_country'][wish['origin_country'] == ''] <- 'Not defined'

#shipping_is_express has too many zero, so we decided to exclude this column
wish <- select(wish, -c(shipping_is_express))

#Only 7 colors have more than 100 records so We decided to keep only 8 factors of color, i.e. black, white, blue, red, green, yellow, pink and others.
color_list <- c('black', 'white', 'blue', 'red', 'green', 'yellow', 'pink')
wish$product_color[!(wish$product_color %in% color_list)] <- 'others'

wish %>%
  group_by(product_color) %>%
  summarise(no_rows = length(product_color)) %>%
  arrange(desc(no_rows)) %>%
  filter(no_rows > 100)

#We decided to change origin to CN and others.
wish$origin_country <- as.character(wish$origin_country)
wish$origin_country[which(wish$origin_country != 'CN')] <- 'others'
wish$origin_country[is.na(wish$origin_country)] <- 'others'

wish %>%
  group_by(origin_country) %>%
  summarise(no_rows = length(origin_country)) %>%
  arrange(desc(no_rows))

#convert column name to short version
origin_colname <- colnames(wish)
colnames(wish) <- c('price', 'retail', 'sold_ct', 'ad_boost', 'rate', 'rate_ct', 'rate5', 'rate4', 'rate3', 'rate2', 'rate1', 'badge_ct', 'bg_local', 'bg_quality', 'bg_fastship', 'color', 'size', 'inventory', 'ship_price', 'able_country', 'total_invent', 'has_bg_urgency', 'origin', 'seller_rate_ct', 'seller_rate', 'has_seller_propic')


library(corrplot)

## corrplot 0.84 loaded

# finding correlation between numeric columns and charges

numeric.column <- sapply(wish, is.numeric)
corr <- cor(wish[, numeric.column]) #, use = 'pairwise.complete.obs'
corrplot(corr, method = 'color')

```



#convert the y (sold_ct) to categorical. Also since it is unbalanced we group some category together.

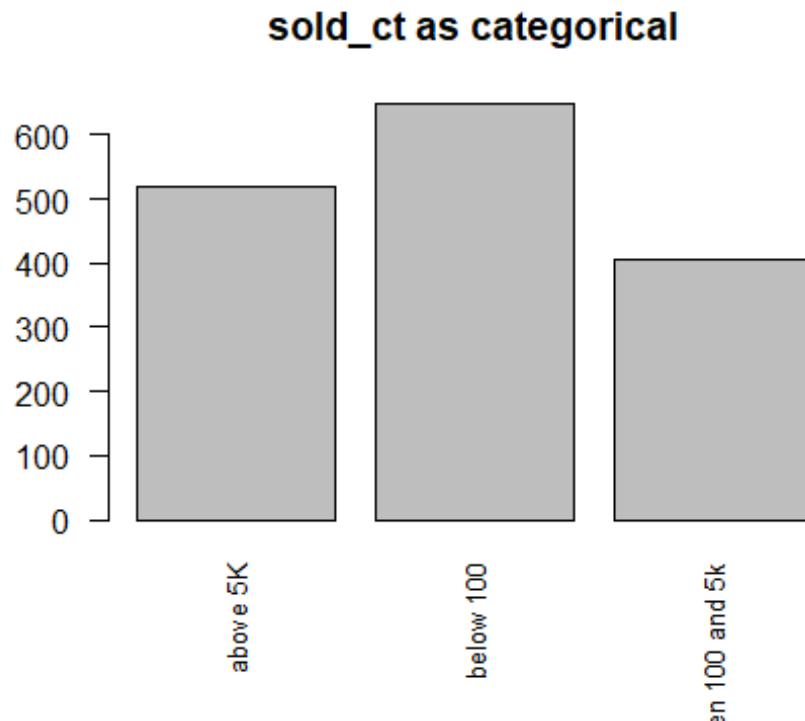
```
table(wish['sold_ct']) # very unbalanced
```

```
##
##      1      2      3      6      7      8     10     50    100   1000   50
00
##      3      2      2      1      2      4     49     76    509   405    2
17
## 10000 20000 50000 100000
##   177   103    17     6
```

```
wish_cate <- wish
wish_cate$sold_ct_cate <- wish_cate$sold_ct
wish_cate$sold_ct_cate[which(wish_cate$sold_ct <= 100)] <- 'below 100'
wish_cate$sold_ct_cate[which(wish_cate$sold_ct >= 5000)] <- 'above 5K'
wish_cate$sold_ct_cate[which(wish_cate$sold_ct > 100 & wish_cate$sold_ct < 5000)] <- 'between 100 and 5k'
wish_cate <- select(wish_cate, -sold_ct)
wish_cate$sold_ct_cate <- as.factor(wish_cate$sold_ct_cate)
wish_cate$color <- as.factor(wish_cate$color)
wish_cate$size <- as.factor(wish_cate$size)
wish_cate$origin <- as.factor(wish_cate$origin)
table(wish_cate$sold_ct_cate) # much better
```

```
##
##           above 5K           below 100  between 100 and 5k
##           520                648                405

x1 <- factor(wish_cate$sold_ct_cate)
tb <- table(x1)
barplot(tb, names.arg = row.names(tb), cex.names = 0.8, main = "sold_ct as ca
tegorical", las = 2)
```



```
#percentage of each rate count
wish_cate$rate5_pct <- wish_cate$rate5/wish_cate$rate_ct
wish_cate$rate4_pct <- wish_cate$rate4/wish_cate$rate_ct
wish_cate$rate3_pct <- wish_cate$rate3/wish_cate$rate_ct
wish_cate$rate2_pct <- wish_cate$rate2/wish_cate$rate_ct
wish_cate$rate1_pct <- wish_cate$rate1/wish_cate$rate_ct

drops <- c('rate_ct', 'rate5', 'rate4', 'rate3', 'rate2', 'rate1')
wish_cate <- wish_cate[, !(names(wish_cate) %in% drops)]

wish_cate <- wish_cate %>% drop_na(rate5_pct)
wish_cate <- wish_cate %>% drop_na(price)

summary(wish_cate)

##           price           retail           ad_boost           rate
##  Min.   : 1.000   Min.   : 1.00   Min.   :0.0000   Min.   :1.000
##  1st Qu.: 5.830   1st Qu.: 7.00   1st Qu.:0.0000   1st Qu.:3.530
```

```

## Median : 8.000 Median : 10.00 Median :0.0000 Median :3.830
## Mean : 8.335 Mean : 23.27 Mean :0.4332 Mean :3.786
## 3rd Qu.:11.000 3rd Qu.: 26.00 3rd Qu.:1.0000 3rd Qu.:4.090
## Max. :49.000 Max. :252.00 Max. :1.0000 Max. :5.000
##
## badge_ct bg_local bg_quality bg_fastship
## Min. :0.0000 Min. :0.00000 Min. :0.00000 Min. :0.00000
## 1st Qu.:0.0000 1st Qu.:0.00000 1st Qu.:0.00000 1st Qu.:0.00000
## Median :0.0000 Median :0.00000 Median :0.00000 Median :0.00000
## Mean :0.1086 Mean :0.01898 Mean :0.07657 Mean :0.01309
## 3rd Qu.:0.0000 3rd Qu.:0.00000 3rd Qu.:0.00000 3rd Qu.:0.00000
## Max. :3.0000 Max. :1.00000 Max. :1.00000 Max. :1.00000
##
## color size inventory ship_price
## black :304 l : 53 Min. : 1.00 Min. : 1.000
## others :278 m :204 1st Qu.: 6.00 1st Qu.: 2.000
## white :268 One-sized: 87 Median :50.00 Median : 2.000
## blue :167 s :669 Mean :33.22 Mean : 2.345
## red :150 xl : 56 3rd Qu.:50.00 3rd Qu.: 3.000
## green :138 xs :459 Max. :50.00 Max. :12.000
## (Other):223
## able_country total_invent has_bg_urgency origin
## Min. : 6.00 Min. : 1.00 Min. :0.000 CN :1472
## 1st Qu.: 31.00 1st Qu.:50.00 1st Qu.:0.000 others: 56
## Median : 40.00 Median :50.00 Median :0.000
## Mean : 40.45 Mean :49.82 Mean :0.301
## 3rd Qu.: 43.00 3rd Qu.:50.00 3rd Qu.:1.000
## Max. :140.00 Max. :50.00 Max. :1.000
##
## seller_rate_ct seller_rate has_seller_propic sold_ct_c
ate
## Min. : 3 Min. :2.941 Min. :0.0000 above 5K :52
0
## 1st Qu.: 2116 1st Qu.:3.919 1st Qu.:0.0000 below 100 :60
3
## Median : 8194 Median :4.041 Median :0.0000 between 100 and 5k:40
5
## Mean : 26667 Mean :4.033 Mean :0.1466
## 3rd Qu.: 24616 3rd Qu.:4.160 3rd Qu.:0.0000
## Max. :2174765 Max. :4.578 Max. :1.0000
##
## rate5_pct rate4_pct rate3_pct rate2_pct
## Min. :0.0000 Min. :0.0000 Min. :0.0000 Min. :0.00000
## 1st Qu.:0.3889 1st Qu.:0.1603 1st Qu.:0.1082 1st Qu.:0.04147
## Median :0.4715 Median :0.1879 Median :0.1429 Median :0.06667
## Mean :0.4680 Mean :0.1861 Mean :0.1406 Mean :0.07467
## 3rd Qu.:0.5537 3rd Qu.:0.2143 3rd Qu.:0.1693 3rd Qu.:0.09140
## Max. :1.0000 Max. :0.6667 Max. :0.5000 Max. :1.00000
##
## rate1_pct

```

```
## Min.      :0.00000
## 1st Qu.:0.07021
## Median :0.11073
## Mean     :0.13069
## 3rd Qu.:0.16874
## Max.      :1.00000
##

str(wish_cate)

## 'data.frame':    1528 obs. of  25 variables:
## $ price      : num  16 8 8 8 2.72 3.92 7 12 11 5.78 ...
## $ retail     : int  14 22 43 8 3 9 6 11 84 22 ...
## $ ad_boost   : int  0 1 0 1 1 0 0 0 1 0 ...
## $ rate       : num  3.76 3.45 3.57 4.03 3.1 5 3.84 3.76 3.47 3.6 ..
.
## $ badge_ct   : int  0 0 0 0 0 0 0 0 0 0 ...
## $ bg_local   : int  0 0 0 0 0 0 0 0 0 0 ...
## $ bg_quality : int  0 0 0 0 0 0 0 0 0 0 ...
## $ bg_fastship : int  0 0 0 0 0 0 0 0 0 0 ...
## $ color      : Factor w/ 8 levels "black","blue",...: 7 3 8 1 8 2 7
2 1 4 ...
## $ size       : Factor w/ 6 levels "l","m","One-sized",...: 2 6 6 2 4
6 6 2 2 4 ...
## $ inventory  : int  50 50 1 50 1 1 50 50 50 50 ...
## $ ship_price : int  4 2 3 2 1 1 2 3 2 2 ...
## $ able_country : int  34 41 36 41 35 40 31 139 36 33 ...
## $ total_invent : int  50 50 50 50 50 50 50 50 50 50 ...
## $ has_bg_urgency : num  1 1 1 0 1 0 0 0 1 0 ...
## $ origin     : Factor w/ 2 levels "CN","others": 1 1 1 1 1 1 1 1 1
1 ...
## $ seller_rate_ct : int  568 17752 295 23832 14482 65 10194 342 330 5534
...
## $ seller_rate   : num  4.13 3.9 3.99 4.02 4 ...
## $ has_seller_propic: int  0 0 0 0 0 0 1 0 0 0 ...
## $ sold_ct_cate  : Factor w/ 3 levels "above 5K","below 100",...: 2 1 2
1 2 2 1 3 2 1 ...
## $ rate5_pct     : num  0.481 0.37 0.357 0.509 0.3 ...
## $ rate4_pct     : num  0.148 0.167 0.286 0.206 0.2 ...
## $ rate3_pct     : num  0.185 0.182 0.143 0.15 0.1 ...
## $ rate2_pct     : num  0.0185 0.105 0 0.0725 0.1 ...
## $ rate1_pct     : num  0.1667 0.1756 0.2143 0.0622 0.3 ...
```

Methodology

80:20 split for train and test set

```
set.seed(123)

train_rows <- sample(1:nrow(wish), 0.8 * nrow(wish))
```

```
wish.train <- wish_cate[train_rows, ] # wish training set
wish.test <- wish_cate[-train_rows, ]

wish.train <- wish.train %>% drop_na(price)
```

Multinomial Regression

```
set.seed(123)
library(nnet)
multinomial.mod <- multinom(sold_ct_cate ~ ., data = wish.train) #, na.action
= na.roughfix

## # weights: 108 (70 variable)
## initial value 1339.208380
## iter 10 value 1255.977483
## iter 20 value 1189.343320
## iter 30 value 1164.442293
## iter 40 value 1161.648181
## iter 50 value 1160.944105
## iter 60 value 1160.541882
## iter 70 value 1159.267178
## iter 80 value 1158.778710
## final value 1158.151626
## converged

summary(multinomial.mod)

## Call:
## multinom(formula = sold_ct_cate ~ ., data = wish.train)
##
## Coefficients:
## (Intercept) price retail ad_boost rate
## below 100 61.52575 0.05993473 0.007125439 0.02511886 52.47417
## between 100 and 5k 162.13209 0.05669671 0.006771135 0.19507681 8.30567
## badge_ct bg_local bg_quality bg_fastship colorblue
## below 100 -0.09765252 -1.0448405 -0.07811868 1.02530663 -0.17996
180
## between 100 and 5k 0.01625992 -0.2818895 0.24682887 0.05132056 -0.03807
248
## colorgreen colorothers colorpink colorred colorwhite
## below 100 0.2602755 -0.187931448 -0.0856353 0.1882070 -0.2789398
1
## between 100 and 5k 0.3371423 -0.000843693 0.4433978 0.6170049 -0.0426730
5
## coloryellow sizem sizeOne-sized sizes size
exl
## below 100 0.9822648 0.1970586 1.5743267 0.41382825 0.8979
739
```



```

## between 100 and 5k 1.2630431 -0.3892427 0.5845379 0.01774944 -0.3659
442
##          sizexs    inventory  ship_price able_country total_i
nvent
## below 100      1.3233542 -0.011385347 -0.17951453 0.01114518 -4.4
93798
## between 100 and 5k 0.5343495 -0.006466085 -0.05803779 0.01434421 -4.3
65661
##          has_bg_urgency originothers seller_rate_ct seller_rate
## below 100      0.09241284 1.0077496 -1.677742e-05 -1.6910850
## between 100 and 5k 0.17423052 0.1463904 -1.292792e-05 -0.5705007
##          has_seller_propic rate5_pct rate4_pct rate3_pct rate2_p
ct
## below 100      -0.33584439 -93.12020 -40.06576 8.845931 65.005
29
## between 100 and 5k 0.07444447 15.54658 24.04949 32.023675 39.895
97
##          rate1_pct
## below 100      120.86049
## between 100 and 5k 50.61638
##
## Std. Errors:
##          (Intercept)      price      retail      ad_boost
## below 100      8.607617e-05 0.0005094232 0.002561476 5.506705e-05
## between 100 and 5k 9.088986e-05 0.0006262238 0.002580111 6.270339e-05
##          rate      badge_ct      bg_local      bg_quality
## below 100      0.0003177981 4.290554e-06 1.452266e-06 5.115773e-06
## between 100 and 5k 0.0003372213 9.832451e-06 3.080471e-06 8.086474e-06
##          bg_fastship colorblue colorgreen colorothers
## below 100      1.197296e-06 6.700695e-06 1.170119e-05 2.473794e-05
## between 100 and 5k 8.439712e-07 6.012753e-06 1.348777e-05 2.312286e-05
##          colorpink colorred colorwhite coloryellow
## below 100      7.596063e-06 1.558492e-05 2.029067e-05 1.036194e-05
## between 100 and 5k 8.749608e-06 1.770766e-05 1.659566e-05 1.452224e-05
##          sizem sizeOne-sized      sizes      sizexl
## below 100      9.245807e-06 4.306001e-06 2.445879e-05 2.957718e-06
## between 100 and 5k 4.759126e-06 6.895911e-06 3.000295e-05 1.796344e-06
##          sizexs    inventory  ship_price able_country
## below 100      8.797572e-05 0.003763135 0.0001556580 0.004116648
## between 100 and 5k 8.126724e-05 0.004002453 0.0001837391 0.004186555
##          total_invent has_bg_urgency originothers seller_rate_ct
## below 100      0.004305804 2.088730e-05 2.187938e-06 2.664673e-06
## between 100 and 5k 0.004542852 2.247189e-05 1.212141e-06 2.407260e-06
##          seller_rate has_seller_propic      rate5_pct      rate4_pc
t
## below 100      0.0003359766      1.283221e-05 3.756664e-05 1.633464e-0
5
## between 100 and 5k 0.0003572466      2.017580e-05 3.987934e-05 1.703798e-0
5
##          rate3_pct      rate2_pct      rate1_pct

```

```

## below 100          1.258024e-05 7.464148e-06 1.229044e-05
## between 100 and 5k 1.458081e-05 6.645865e-06 1.281971e-05
##
## Residual Deviance: 2316.303
## AIC: 2448.303

multinomial.pred_train <- predict(multinomial.mod, wish.train)
multinomial.pred_test <- predict(multinomial.mod, wish.test)
# training error
print("Misclassification rate on the training set:")

## [1] "Misclassification rate on the training set:"

mean(as.character(multinomial.pred_train) != as.character(wish.train$sold_ct_
cate))

## [1] 0.4577523

# test error
print("Misclassification rate on the test set:")

## [1] "Misclassification rate on the test set:"

mean(as.character(multinomial.pred_test) != as.character(wish.test$sold_ct_ca
te))

## [1] 0.4854369

confusion.matrix <- table(wish.test$sold_ct_cate, multinomial.pred_test)
print(confusion.matrix)

##
##          multinomial.pred_test
##          above 5K below 100 between 100 and 5k
## above 5K          61      28          15
## below 100         27      90           7
## between 100 and 5k 26      47           8

accuracy.percent <- 100*sum(diag(confusion.matrix))/sum(confusion.matrix)
above5k.precent <- 100*confusion.matrix[1,1]/sum(confusion.matrix[1,])
print(paste("Test accuracy:",accuracy.percent,"%"))

## [1] "Test accuracy: 51.4563106796116 %"

print(paste("Above 5k accuracy:",above5k.precent,"%"))

## [1] "Above 5k accuracy: 58.6538461538462 %"

```

Dicision Tree Models

```
library(tree)
```

```

## Registered S3 method overwritten by 'tree':
##   method      from
##   print.tree cli

set.seed(123)

tree.wish <- tree(sold_ct_cate ~ ., data = wish.train)
summary(tree.wish)

##
## Classification tree:
## tree(formula = sold_ct_cate ~ ., data = wish.train)
## Variables actually used in tree construction:
## [1] "rate2_pct"      "seller_rate_ct" "rate3_pct"      "rate4_pct"
## Number of terminal nodes: 10
## Residual mean deviance: 1.418 = 1715 / 1209
## Misclassification error rate: 0.3158 = 385 / 1219

tree.pred <- predict(tree.wish, wish.test, type = "class")

# table(tree.pred, wish.test$sold_ct_cate)

# print("Misclassification error rate on test set: ")

confusion.matrix <- table(wish.test$sold_ct_cate, tree.pred)
print(confusion.matrix)

##
##               tree.pred
##               above 5K below 100 between 100 and 5k
## above 5K           88      5              11
## below 100          13     98              13
## between 100 and 5k  42     20              19

accuracy.percent <- 100*sum(diag(confusion.matrix))/sum(confusion.matrix)
above5k.precent <- 100*confusion.matrix[1,1]/sum(confusion.matrix[1,])
print(paste("Test accuracy:",accuracy.percent,"%"))

## [1] "Test accuracy: 66.3430420711974 %"

print(paste("Above 5k accuracy:",above5k.precent,"%"))

## [1] "Above 5k accuracy: 84.6153846153846 %"

```

Bagging

```

library(randomForest)

## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.

```

```
##
## Attaching package: 'randomForest'

## The following object is masked from 'package:ggplot2':
##
##      margin

## The following object is masked from 'package:dplyr':
##
##      combine

set.seed(123)

bag.wish <- randomForest(sold_ct_cate ~ ., data = wish.train, mtry = length(wish.train) - 1, importance = TRUE, na.action = na.roughfix)
bag.wish

##
## Call:
## randomForest(formula = sold_ct_cate ~ ., data = wish.train, mtry = length(wish.train) - 1, importance = TRUE, na.action = na.roughfix)
##
##      Type of random forest: classification
##      Number of trees: 500
## No. of variables tried at each split: 24
##
##      OOB estimate of  error rate: 27.07%
## Confusion matrix:
##
##      above 5K below 100 between 100 and 5k class.error
## above 5K      329      10          77    0.2091346
## below 100       20     419          40    0.1252610
## between 100 and 5k  122      61         141    0.5648148

bag.pred <- predict(bag.wish, wish.test)

confusion.matrix <- table(wish.test$sold_ct_cate, bag.pred)
print(confusion.matrix)

##
##      bag.pred
##      above 5K below 100 between 100 and 5k
## above 5K      84       3          17
## below 100       3     111          10
## between 100 and 5k  24      12          45

accuracy.percent <- 100*sum(diag(confusion.matrix))/sum(confusion.matrix)
above5k.precent <- 100*confusion.matrix[1,1]/sum(confusion.matrix[1,])
print(paste("Test accuracy:",accuracy.percent,"%"))

## [1] "Test accuracy: 77.6699029126214 %"

print(paste("Above 5k accuracy:",above5k.precent,"%"))

## [1] "Above 5k accuracy: 80.7692307692308 %"
```

Random forest

```
set.seed(123)

rf.wish <- randomForest(sold_ct_cate ~ ., data = wish.train, mtry = (length(wish.train) - 1) / 3, importance = TRUE, na.action = na.roughfix)
rf.wish

##
## Call:
## randomForest(formula = sold_ct_cate ~ ., data = wish.train, mtry = (length(wish.train) - 1)/3, importance = TRUE, na.action = na.roughfix)
##
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 8
##
##           OOB estimate of  error rate: 27.24%
## Confusion matrix:
##
##           above 5K below 100 between 100 and 5k class.error
## above 5K           331         13             72    0.2043269
## below 100           23        419             37    0.1252610
## between 100 and 5k   122         65            137    0.5771605

rf.pred <- predict(rf.wish, wish.test)

confusion.matrix <- table(wish.test$sold_ct_cate, rf.pred)
print(confusion.matrix)

##
##           rf.pred
##           above 5K below 100 between 100 and 5k
## above 5K           83         2             19
## below 100           4        108             12
## between 100 and 5k   23         12             46

accuracy.percent <- 100*sum(diag(confusion.matrix))/sum(confusion.matrix)
above5k.precent <- 100*confusion.matrix[1,1]/sum(confusion.matrix[1,])
print(paste("Test accuracy:",accuracy.percent,"%"))

## [1] "Test accuracy: 76.6990291262136 %"

print(paste("Above 5k accuracy:",above5k.precent,"%"))

## [1] "Above 5k accuracy: 79.8076923076923 %"
```

SVM

Linear

```
library(e1071)
# summary(wish.train)
```

```

# summary(wish.test)

set.seed(123)

tuned <- tune(svm, sold_ct_cate ~ ., data = wish.train, kernel = "linear", ranges = list(cost = append(seq(0.01, 10, by = 0.5), 10)))
summary(tuned)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   8.51
##
## - best performance: 0.4814862
##
## - Detailed performance results:
##   cost      error dispersion
## 1  0.01 0.5053109 0.04242946
## 2  0.51 0.4839385 0.04506527
## 3  1.01 0.4839656 0.04133876
## 4  1.51 0.4839724 0.04068346
## 5  2.01 0.4814998 0.03790904
## 6  2.51 0.4815066 0.03637876
## 7  3.01 0.4814998 0.04201858
## 8  3.51 0.4831459 0.03843672
## 9  4.01 0.4831391 0.04081346
## 10 4.51 0.4839588 0.04090609
## 11 5.01 0.4847785 0.04098030
## 12 5.51 0.4831324 0.04218639
## 13 6.01 0.4847853 0.03975657
## 14 6.51 0.4839520 0.04349467
## 15 7.01 0.4831256 0.04569449
## 16 7.51 0.4831324 0.04340741
## 17 8.01 0.4839520 0.04434453
## 18 8.51 0.4814862 0.04346441
## 19 9.01 0.4823127 0.04347480
## 20 9.51 0.4814862 0.04380657
## 21 10.00 0.4831324 0.04357905

print("The best cost:")

## [1] "The best cost:"

tuned$best.parameter$cost

## [1] 8.51

```

```

lin.svm <- svm(sold_ct_cate ~ ., kernel = "linear", type = "C-class", data =
wish.train, cost = tuned$best.parameter$cost)

train_pred <- predict(lin.svm, wish.train, na.action = na.exclude)
table <- table(wish.train$sold_ct_cate, train_pred)

print(paste("accuracy with cost =", tuned$best.parameter$cost, "for train: "))
)

## [1] "accuracy with cost = 8.51 for train: "
1-(sum(table)-sum(diag(table))) / (sum(table))

## [1] 0.5348646

test_pred <- predict(lin.svm, wish.test)
table <- table(wish.test$sold_ct_cate, test_pred)

print(paste("accuracy with cost =", tuned$best.parameter$cost, "for test: "))

## [1] "accuracy with cost = 8.51 for test: "
1-(sum(table)-sum(diag(table))) / (sum(table))

## [1] 0.5048544

print(paste("above 5k group - accuracy with cost =", tuned$best.parameter$cost,
": "))

## [1] "above 5k group - accuracy with cost = 8.51 : "
sum(table[1,1]) / sum(table[1,])

## [1] 0.5384615

# we cannot plot SVM classification plot since we have more than 2 columns
table(wish.test$sold_ct_cate, test_pred)

##
##          test_pred
##          above 5K below 100 between 100 and 5k
##  above 5K          56      35          13
##  below 100         18      91          15
##  between 100 and 5k  26      46           9

```

Radial

```

# names(wish.train)
set.seed(123)

tuned <- tune(svm, sold_ct_cate ~ ., data = wish.train, kernel = "radial", ranges = list(cost = append(seq(0.01, 15, by = 0.5), 10)))
summary(tuned)

```

```

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   12.01
##
## - best performance: 0.3675112
##
## - Detailed performance results:
##   cost      error dispersion
## 1  0.01 0.6070519 0.04260201
## 2  0.51 0.4298537 0.03659364
## 3  1.01 0.4159193 0.03849578
## 4  1.51 0.4077293 0.03480182
## 5  2.01 0.3995258 0.03606130
## 6  2.51 0.3970532 0.03736618
## 7  3.01 0.3904891 0.04562588
## 8  3.51 0.3921284 0.04731593
## 9  4.01 0.3863907 0.04793160
## 10 4.51 0.3822924 0.03879900
## 11 5.01 0.3814862 0.03992683
## 12 5.51 0.3823059 0.03877121
## 13 6.01 0.3814795 0.03870461
## 14 6.51 0.3781940 0.03609755
## 15 7.01 0.3790137 0.03875255
## 16 7.51 0.3814659 0.03912204
## 17 8.01 0.3839249 0.03433190
## 18 8.51 0.3822788 0.03786632
## 19 9.01 0.3781872 0.03880695
## 20 9.51 0.3749085 0.03738736
## 21 10.01 0.3757282 0.03752909
## 22 10.51 0.3699837 0.03683100
## 23 11.01 0.3683376 0.03636024
## 24 11.51 0.3683308 0.03789495
## 25 12.01 0.3675112 0.03489759
## 26 12.51 0.3691505 0.03202574
## 27 13.01 0.3716095 0.03338185
## 28 13.51 0.3732557 0.03371330
## 29 14.01 0.3691573 0.03347686
## 30 14.51 0.3675180 0.03184827
## 31 10.00 0.3749085 0.03718716

print("The best cost:")
## [1] "The best cost:"
tuned$best.parameter$cost

```



```
## [1] 12.01

rad.svm <- svm(sold_ct_cate ~ ., kernel = "radial", data = wish.train, cost =
tuned$best.parameter$cost)

train_pred <- predict(rad.svm, wish.train, na.action = na.exclude)
table <- table(wish.train$sold_ct_cate, train_pred)

print(paste("accuracy with cost =", tuned$best.parameter$cost, "for train: ")
)

## [1] "accuracy with cost = 12.01 for train: "
1-(sum(table)-sum(diag(table))) / (sum(table))

## [1] 0.8630025

test_pred <- predict(rad.svm, wish.test)
table <- table(wish.test$sold_ct_cate, test_pred)

print(paste("accuracy with cost =", tuned$best.parameter$cost, "for test: "))

## [1] "accuracy with cost = 12.01 for test: "
1-(sum(table)-sum(diag(table))) / (sum(table))

## [1] 0.6343042

print(paste("above 5k group - accuracy with cost =", tuned$best.parameter$cost,
": "))

## [1] "above 5k group - accuracy with cost = 12.01 : "
sum(table[1,1]) / sum(table[1,])

## [1] 0.6730769

table(wish.test$sold_ct_cate, test_pred)

##               test_pred
##               above 5K below 100 between 100 and 5k
##   above 5K           70      8           26
##   below 100          11     93           20
##   between 100 and 5k  36     12           33
```

The result shows that there is overfitting issue. (training error is getting low, but test error is getting higher)

Polynomial

```
set.seed(123)
tune.poly <- tune(svm, sold_ct_cate ~ ., data = wish.train, kernel = "poly",
```

```
degree = 3, ranges = list(cost = append(seq(0.01, 15, by = 0.5), 10)))
summary(tuned)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
## 12.01
##
## - best performance: 0.3675112
##
## - Detailed performance results:
##   cost      error dispersion
## 1  0.01 0.6070519 0.04260201
## 2  0.51 0.4298537 0.03659364
## 3  1.01 0.4159193 0.03849578
## 4  1.51 0.4077293 0.03480182
## 5  2.01 0.3995258 0.03606130
## 6  2.51 0.3970532 0.03736618
## 7  3.01 0.3904891 0.04562588
## 8  3.51 0.3921284 0.04731593
## 9  4.01 0.3863907 0.04793160
## 10 4.51 0.3822924 0.03879900
## 11 5.01 0.3814862 0.03992683
## 12 5.51 0.3823059 0.03877121
## 13 6.01 0.3814795 0.03870461
## 14 6.51 0.3781940 0.03609755
## 15 7.01 0.3790137 0.03875255
## 16 7.51 0.3814659 0.03912204
## 17 8.01 0.3839249 0.03433190
## 18 8.51 0.3822788 0.03786632
## 19 9.01 0.3781872 0.03880695
## 20 9.51 0.3749085 0.03738736
## 21 10.01 0.3757282 0.03752909
## 22 10.51 0.3699837 0.03683100
## 23 11.01 0.3683376 0.03636024
## 24 11.51 0.3683308 0.03789495
## 25 12.01 0.3675112 0.03489759
## 26 12.51 0.3691505 0.03202574
## 27 13.01 0.3716095 0.03338185
## 28 13.51 0.3732557 0.03371330
## 29 14.01 0.3691573 0.03347686
## 30 14.51 0.3675180 0.03184827
## 31 10.00 0.3749085 0.03718716
```

```
print("The best cost:")
```

```

## [1] "The best cost:"
tuned$best.parameter$cost

## [1] 12.01

poly.svm <- svm(sold_ct_cate ~ ., kernel = "poly", data = wish.train, degree
= 3, cost = tuned$best.parameter$cost)

train_pred <- predict(poly.svm, wish.train, na.action = na.exclude)
table <- table(wish.train$sold_ct_cate, train_pred)

print(paste("accuracy with cost =", tuned$best.parameter$cost, "for train: "))
)

## [1] "accuracy with cost = 12.01 for train: "
1-(sum(table)-sum(diag(table))) / (sum(table))

## [1] 0.733388

test_pred <- predict(poly.svm, wish.test)
table <- table(wish.test$sold_ct_cate, test_pred)

print(paste("accuracy with cost =", tuned$best.parameter$cost, "for test: "))

## [1] "accuracy with cost = 12.01 for test: "
1-(sum(table)-sum(diag(table))) / (sum(table))

## [1] 0.5566343

print(paste("above 5k group - accuracy with cost =", tuned$best.parameter$cost, ": "))

## [1] "above 5k group - accuracy with cost = 12.01 : "
sum(table[1,1]) / sum(table[1,])

## [1] 0.8461538

table(wish.test$sold_ct_cate, test_pred)

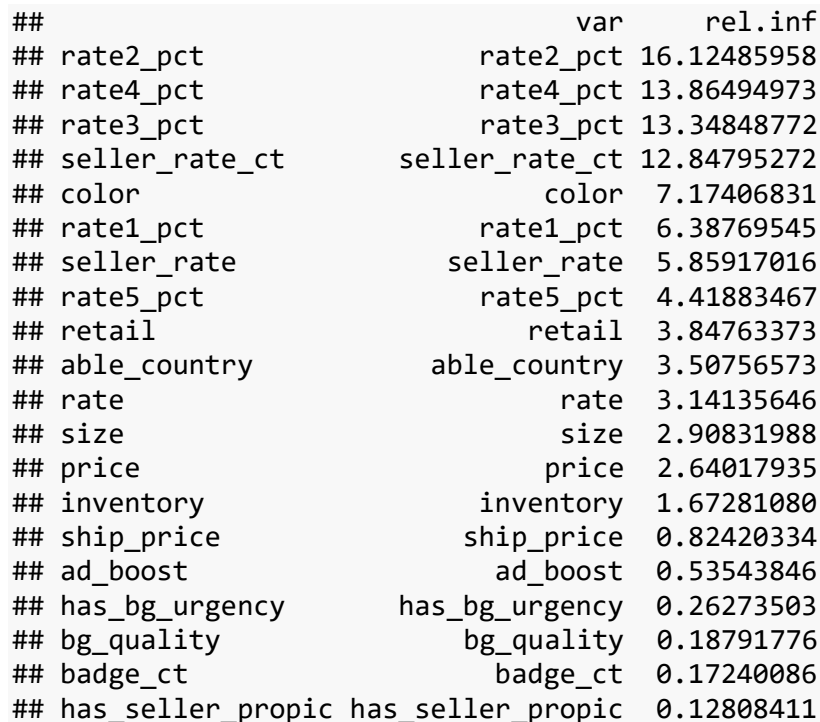
##
##           test_pred
##           above 5K below 100 between 100 and 5k
##   above 5K           88         5              11
##   below 100          35        74              15
##   between 100 and 5k  59        12              10

```

GBM

```
library(gbm)
```

```
boost.wish = gbm(sold_ct_cate ~ ., data = wish.train, distribution = "multino  
mial", n.trees = 10000, shrinkage = 0.01, interaction.depth = 4)  
summary(boost.wish)
```



```
## origin                                origin  0.08833278
## bg_local                             bg_local  0.05700338
## bg_fastship                          bg_fastship 0.00000000
## total_invent                         total_invent 0.00000000

boost.predP <- predict(boost.wish, wish.test, n.trees = 10000, type = 'response')

classification <- c("above 5K", "below 100", "between 100 and 5k")
boost.pred <- apply(boost.predP, 1, which.max)
boost.pred <- classification[boost.pred]
confusion.matrix <- table(wish.test$sold_ct_cate, boost.pred)
print(confusion.matrix)

##                boost.pred
##                above 5K below 100 between 100 and 5k
##  above 5K                82         3                19
##  below 100                4        108                12
##  between 100 and 5k       20        10                51

accuracy.percent <- 100*sum(diag(confusion.matrix))/sum(confusion.matrix)
above5k.precent <- 100*confusion.matrix[1,1]/sum(confusion.matrix[1,])
print(paste("Test accuracy:", accuracy.percent, "%"))

## [1] "Test accuracy: 77.9935275080906 %"

print(paste("Above 5k accuracy:", above5k.precent, "%"))

## [1] "Above 5k accuracy: 78.8461538461538 %"
```

XGBoost

```
library(xgboost)
# Create numeric labels with one-hot encoding
set.seed(123)
train_labs <- as.numeric(wish.train$sold_ct_cate) - 1
val_labs <- as.numeric(wish.test$sold_ct_cate) - 1

# options(na.action='na.pass')
new_train <- model.matrix(~ . + 0, data = subset(wish.train, select = -sold_ct_cate))
new_val <- model.matrix(~ . + 0, data = subset(wish.test, select = -sold_ct_cate))

# Prepare matrices
xgb_train <- xgb.DMatrix(data = new_train, label = train_labs)
xgb_val <- xgb.DMatrix(data = new_val, label = val_labs)

params <- list(booster = "gbtree", objective = "multi:softprob", num_class = 4, eval_metric = "mlogloss")
```

```
# Calculate # of folds for cross-validation
```

```
xgbcv <- xgb.cv(params = params, data = xgb_train, nrounds = 100, nfold = 5,  
showsd = TRUE, stratified = TRUE, print_every_n = 10, early_stop_round = 20,  
maximize = FALSE, prediction = TRUE)
```

```
## [1] train-mlogloss:1.128172+0.005632 test-mlogloss:1.169920+0.009762  
## [11] train-mlogloss:0.433942+0.010224 test-mlogloss:0.684675+0.038885  
## [21] train-mlogloss:0.230444+0.008175 test-mlogloss:0.632566+0.043106  
## [31] train-mlogloss:0.140213+0.007266 test-mlogloss:0.634191+0.042982  
## [41] train-mlogloss:0.088632+0.005966 test-mlogloss:0.638861+0.046647  
## [51] train-mlogloss:0.060192+0.003827 test-mlogloss:0.654579+0.053184  
## [61] train-mlogloss:0.042890+0.002283 test-mlogloss:0.673031+0.057878  
## [71] train-mlogloss:0.032716+0.001623 test-mlogloss:0.691112+0.056010  
## [81] train-mlogloss:0.025648+0.001521 test-mlogloss:0.707691+0.058055  
## [91] train-mlogloss:0.020897+0.001250 test-mlogloss:0.723342+0.062724  
## [100] train-mlogloss:0.017756+0.001176 test-mlogloss:0.734206+0.0604  
95
```

```
# Function to compute classification error
```

```
classification_error <- function(conf_mat) {  
  conf_mat = as.matrix(conf_mat)  
  
  error = 1 - sum(diag(conf_mat)) / sum(conf_mat)  
  
  return (error)  
}
```

```
# Mutate xgb output to deliver hard predictions
```

```
xgb_train_preds <- data.frame(xgbcv$pred) %>% mutate(max = max.col(., ties.me  
thod = "last"), label = train_labs + 1)
```

```
# Examine output
```

```
head(xgb_train_preds)
```

```
##           X1           X2           X3           X4 max label  
## 1 2.557963e-01 0.394735545 3.491999e-01 2.682663e-04 2      3  
## 2 9.978531e-01 0.001962732 1.680852e-04 1.610581e-05 1      1  
## 3 8.626200e-05 0.999510407 4.009536e-04 2.427827e-06 2      2  
## 4 1.106577e-04 0.997806489 2.068996e-03 1.382771e-05 2      2  
## 5 3.321923e-02 0.934673190 3.199683e-02 1.107415e-04 2      3  
## 6 8.847255e-06 0.999940276 5.025148e-05 6.239600e-07 2      2
```

```
xgb_conf_mat <- table(true = train_labs + 1, pred = xgb_train_preds$max)
```

```
# Error
```

```
cat("XGB Training Classification Error Rate:", classification_error(xgb_conf_  
mat), "\n")
```

```
## XGB Training Classification Error Rate: 0.2657916
```

```

# predicting / testing on test dataset
xgb_model <- xgb.train(params = params, data = xgb_train, nrounds = 100)

# Predict for validation set
xgb_val_preds <- predict(xgb_model, newdata = xgb_val)

xgb_val_out <- matrix(xgb_val_preds, nrow = 4, ncol = length(xgb_val_preds) /
4) %>%
  t() %>%
  data.frame() %>%
  mutate(max = max.col(., ties.method = "last"), label = val_labels + 1)

# Confusion Matrix
xgb_val_conf <- table(true = val_labels + 1, pred = xgb_val_out$max)

cat("XGB Validation Classification Error Rate:", 1-classification_error(xgb_val_conf), "\n")

## XGB Validation Classification Error Rate: 0.7508091

cat("XGB Validation Classification Error Rate - above 5k:", xgb_val_conf[1,1]
/sum(xgb_val_conf[1,]), "\n")

## XGB Validation Classification Error Rate - above 5k: 0.8076923

```

Stacked Ensembles

```

# we already have gbm.wish for GBM, now build RF model
# https://docs.h2o.ai/h2o/latest-stable/h2o-docs/data-science/stacked-ensembles.html
# Stacked Ensemble model's performance is not so different from those of base learners'
library(h2o)

##
## -----
##
## Your next step is to start H2O:
##   > h2o.init()
##
## For H2O package documentation, ask for help:
##   > ??h2o
##
## After starting H2O, you can use the Web UI at http://localhost:54321
## For more information visit https://docs.h2o.ai
##
## -----
##
## Attaching package: 'h2o'

```

```

## The following objects are masked from 'package:stats':
##
##   cor, sd, var

## The following objects are masked from 'package:base':
##
##   %*%, %in%, &&, ||, apply, as.factor, as.numeric, colnames,
##   colnames<-, ifelse, is.character, is.factor, is.numeric, log,
##   log10, log1p, log2, round, signif, trunc

h2o.init()

## Connection successful!
##
## R is connected to the H2O cluster:
##   H2O cluster uptime:      14 minutes 33 seconds
##   H2O cluster timezone:    America/New_York
##   H2O data parsing timezone: UTC
##   H2O cluster version:     3.32.0.1
##   H2O cluster version age:  2 months and 7 days
##   H2O cluster name:        H2O_started_from_R_maxmo_plv200
##   H2O cluster total nodes:  1
##   H2O cluster total memory: 7.00 GB
##   H2O cluster total cores:  12
##   H2O cluster allowed cores: 12
##   H2O cluster healthy:     TRUE
##   H2O Connection ip:       localhost
##   H2O Connection port:     54321
##   H2O Connection proxy:    NA
##   H2O Internal Security:   FALSE
##   H2O API Extensions:      Amazon S3, Algos, AutoML, Core V3, TargetE
ncoder, Core V4
##   R Version:                R version 3.6.2 (2019-12-12)

wish.train.h2o <- as.h2o(wish.train)

## |
|                                     | 0%
|
|=====| 100%

wish.test.h2o <- as.h2o(wish.test)

## |
|                                     | 0%
|
|=====| 100%

predictors <- c(colnames(wish.train)[1:length(wish.train) - 1])
response <- "sold_ct_cate"

set.seed(123)

```



```

gbm.wish <- h2o.gbm(x = predictors,
                    y = response,
                    nfolds = 5,
                    distribution = "multinomial",
                    keep_cross_validation_predictions = TRUE,
                    training_frame = wish.train.h2o, seed=1)

## |
|                                     | 0%
|=====|                             | 22%
|=====|                             | 100%

rf.wish <- h2o.randomForest(x = predictors,
                            y = response,
                            training_frame = wish.train.h2o,
                            ntrees = 50,
                            nfolds = 5,
                            keep_cross_validation_predictions = TRUE,
                            seed = 1)

## |
|                                     | 0%
|=====|                             | 14%
|=====|                             | 85%
|=====|                             | 100%

ensemble <- h2o.stackedEnsemble(x = predictors,
                                y = response,
                                training_frame = wish.train.h2o,
                                base_models = list(gbm.wish, rf.wish))

## |
|                                     | 0%
|=====|                             | 100%

perf <- h2o.performance(ensemble, newdata = wish.test.h2o)

# Compare to base Learner performance on the test set
perf_gbm_test <- h2o.performance(gbm.wish, newdata = wish.test.h2o)
perf_rf_test <- h2o.performance(rf.wish, newdata = wish.test.h2o)
baselearner_best_auc_test <- max(h2o.auc(perf_gbm_test), h2o.auc(perf_rf_test))
ensemble_auc_test <- h2o.auc(perf)
print(sprintf("Best Base-learner Test AUC:  %s", baselearner_best_auc_test))

```

```

## [1] "Best Base-learner Test AUC:  -Inf"

print(sprintf("Ensemble Test AUC:  %s", ensemble_auc_test))

## character(0)

perf

## H2OMultinomialMetrics: stackedensemble
##
## Test Set Metrics:
## =====
##
## MSE: (Extract with `h2o.mse`) 0.1869437
## RMSE: (Extract with `h2o.rmse`) 0.4323699
## Logloss: (Extract with `h2o.logloss`) 0.5548341
## Mean Per-Class Error: 0.2670902
## Null Deviance: (Extract with `h2o.nulldeviance`) 669.9311
## Residual Deviance: (Extract with `h2o.residual_deviance`) 342.8875
## AIC: (Extract with `h2o.aic`) NaN
## Confusion Matrix: Extract with `h2o.confusionMatrix(<model>, <data>)`
## =====
## Confusion Matrix: Row labels: Actual class; Column labels: Predicted class
##
##          above 5K below 100 between 100 and 5k  Error      Rate
## above 5K          85         2          17 0.1827 = 19 / 104
## below 100         4        107          13 0.1371 = 17 / 124
## between 100 and 5k 29         10          42 0.4815 = 39 / 81
## Totals            118        119          72 0.2427 = 75 / 309
##
## Hit Ratio Table: Extract with `h2o.hit_ratio_table(<model>, <data>)`
## =====
## Top-3 Hit Ratios:
##   k hit_ratio
## 1 1  0.757282
## 2 2  0.961165
## 3 3  1.000000

perf_gbm_test

## H2OMultinomialMetrics: gbm
##
## Test Set Metrics:
## =====
##
## MSE: (Extract with `h2o.mse`) 0.2065348
## RMSE: (Extract with `h2o.rmse`) 0.454461
## Logloss: (Extract with `h2o.logloss`) 0.6037602
## Mean Per-Class Error: 0.2878935
## R^2: (Extract with `h2o.r2`) 0.6518089
## Confusion Matrix: Extract with `h2o.confusionMatrix(<model>, <data>)`
## =====

```

```

## Confusion Matrix: Row labels: Actual class; Column labels: Predicted class
##
##      above 5K  below 100  between 100 and 5k  Error      Rate
## above 5K      87         2          15 0.1635 = 17 / 104
## below 100      3        103          18 0.1694 = 21 / 124
## between 100 and 5k  30         13          38 0.5309 = 43 / 81
## Totals        120        118          71 0.2621 = 81 / 309
##
## Hit Ratio Table: Extract with `h2o.hit_ratio_table(<model>, <data>)`
## =====
## Top-3 Hit Ratios:
##   k hit_ratio
## 1 1  0.737864
## 2 2  0.948220
## 3 3  1.000000

perf_rf_test

## H2OMultinomialMetrics: drf
##
## Test Set Metrics:
## =====
##
## MSE: (Extract with `h2o.mse`) 0.2127525
## RMSE: (Extract with `h2o.rmse`) 0.461251
## Logloss: (Extract with `h2o.logloss`) 0.6123574
## Mean Per-Class Error: 0.3100205
## R^2: (Extract with `h2o.r2`) 0.6413268
## Confusion Matrix: Extract with `h2o.confusionMatrix(<model>, <data>)`
## =====
## Confusion Matrix: Row labels: Actual class; Column labels: Predicted class
##
##      above 5K  below 100  between 100 and 5k  Error      Rate
## above 5K      84         4          16 0.1923 = 20 / 104
## below 100      4        106          14 0.1452 = 18 / 124
## between 100 and 5k  33         15          33 0.5926 = 48 / 81
## Totals        121        125          63 0.2783 = 86 / 309
##
## Hit Ratio Table: Extract with `h2o.hit_ratio_table(<model>, <data>)`
## =====
## Top-3 Hit Ratios:
##   k hit_ratio
## 1 1  0.721683
## 2 2  0.948220
## 3 3  1.000000

# Generate predictions on a test set
pred <- h2o.predict(ensemble, newdata = wish.test.h2o)

## |
|
|
|=====| 100%

```

Neural network

```
library(neuralnet)

##
## Attaching package: 'neuralnet'

## The following object is masked from 'package:dplyr':
##
##      compute

library(nnet)
y_index = grep("sold_ct_cate", names(wish.train))
nn.train = cbind(wish.train[, -y_index], wish.train[, y_index])
names(nn.train)[names(nn.train) == 'wish.train[, y_index]'] <- 'sold_ct_cate'

nn.train = cbind(nn.train[, 1:8], class.ind(as.factor(nn.train$color)), nn.train[, 11:15], nn.train[, 17:24], class.ind(as.factor(nn.train$size)), class.ind(as.factor(nn.train$origin)), class.ind(as.factor(nn.train$sold_ct_cate)))
#normalized by scaling data
scl <- function(x){ (x - min(x))/(max(x) - min(x)) }
nn.train[, 1:29] <- data.frame(lapply(nn.train[, 1:29], scl))

names(nn.train)[names(nn.train) == 'One-sized'] <- 'one_sized'
names(nn.train)[names(nn.train) == 'above 5K'] <- 'above_5k'
names(nn.train)[names(nn.train) == 'below 100'] <- 'below_100'
names(nn.train)[names(nn.train) == 'between 100 and 5k'] <- 'btw_100_5k'

name <- names(nn.train)
f <- as.formula(paste("above_5k + below_100 + btw_100_5k ~", paste(name[!name %in% c("above_5k", "below_100", "btw_100_5k")], collapse = " + ")))
f

## above_5k + below_100 + btw_100_5k ~ price + retail + ad_boost +
##      rate + badge_ct + bg_local + bg_quality + bg_fastship + black +
##      blue + green + others + pink + red + white + yellow + inventory +
##      ship_price + able_country + total_invent + has_bg_urgency +
##      seller_rate_ct + seller_rate + has_seller_propic + rate5_pct +
##      rate4_pct + rate3_pct + rate2_pct + rate1_pct + l + m + one_sized +
##      s + xl + xs + CN + others

set.seed(123)
nn <- neuralnet(f,
  data = nn.train,
  hidden = c(37, 15, 3),
  act.fct = "logistic",
  linear.output = FALSE, stepmax = 3000)

plot(nn)
```

```

nn.test = cbind(wish.test[, -y_index], wish.test[, y_index])
names(nn.test)[names(nn.test) == 'wish.test[, y_index]'] <- 'sold_ct_cate'

nn.test = cbind(nn.test[, 1:8], class.ind(as.factor(nn.test$color)), nn.test[,
11:15], nn.test[, 17:24], class.ind(as.factor(nn.test$size)), class.ind(as.factor(nn.test$origin)), class.ind(as.factor(nn.test$sold_ct_cate)))
#normalized by scaling data
scl <- function(x){ (x - min(x))/(max(x) - min(x)) }
nn.test[, 1:29] <- data.frame(lapply(nn.test[, 1:29], scl))

names(nn.test)[names(nn.test) == 'One-sized'] <- 'one_sized'
names(nn.test)[names(nn.test) == 'above 5K'] <- 'above_5k'
names(nn.test)[names(nn.test) == 'below 100'] <- 'below_100'
names(nn.test)[names(nn.test) == 'between 100 and 5k'] <- 'btw_100_5k'

#train accuracy
nn.train_pred <- compute(nn, nn.train[, 1:37])
nn.train_pred <- nn.train_pred$net.result
true_y.train <- max.col(nn.train[, 38:40])
predicted_y.train <- max.col(nn.train_pred)

#test accuracy
nn.test_pred <- compute(nn, nn.test[, 1:37])
nn.test_pred <- nn.test_pred$net.result
true_y.test <- max.col(nn.test[, 38:40])
predicted_y.test <- max.col(nn.test_pred)

table.train <- table(true_y.train, predicted_y.train)
table.test <- table(true_y.test, predicted_y.test)

print("accuracy for train: ")
## [1] "accuracy for train: "
print(1-(sum(table.train)-sum(diag(table.train))) / (sum(table.train)))
## [1] 0.9589828
print("accuracy for test: ")
## [1] "accuracy for test: "
print(1-(sum(table.test)-sum(diag(table.test))) / (sum(table.test)))
## [1] 0.5210356
print("above 5k group accuracy:")
## [1] "above 5k group accuracy:"

```

```
sum(table.test[1,1]) /sum(table.test[1,])
```

```
## [1] 0.6153846
```